

Topic 3

# Python Deep Learning (DL)

---

AI HUMAN INTERFACE

# Learning Outcomes


---

- Use of Deep Learning (DL) Tools
  - Describe some available python deep learning tools
  - Install and configure Keras & Tensorflow
- Deep Learning architectures
  - Using Model Zoo with pre-trained networks
  - Review of deep learning architectures

# Python Deep Learning

---

Top 20 Python AI and Machine Learning projects on Github



source <https://www.kdnuggets.com/2018/02/top-20-python-ai-machine-learning-open-source-projects.html>

# Open source ML

---

- **TensorFlow** was originally developed by researchers and engineers working on the Google Brain Team within Google's Machine Intelligence research organization. The system is designed to facilitate research in machine learning, and to make it quick and easy to transition from research prototype to production system.
  - <https://github.com/tensorflow/tensorflow>
  - <https://www.tensorflow.org/>
  
- **Keras**, a high-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano.
  - <https://github.com/keras-team/keras>
  - <https://keras.io/>

# Open source ML

---

- ❑ **PyTorch** Tensors and Dynamic neural networks in Python with strong GPU acceleration (No official support for Windows)
  - <https://github.com/pytorch/pytorch>
  - <http://pytorch.org/>
  
- ❑ **CNTK** The Microsoft Cognitive Toolkit (CNTK) is an open-source toolkit for commercial-grade distributed deep learning. It describes neural networks as a series of computational steps via a directed graph.
  - <https://github.com/Microsoft/CNTK>
  - <https://www.microsoft.com/en-us/cognitive-toolkit/>

# TensorFlow

---

## About TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them. The flexible architecture allows you to deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API. TensorFlow was originally developed by researchers and engineers working on the Google Brain Team within Google's Machine Intelligence research organization for the purposes of conducting machine learning and deep neural networks research, but the system is general enough to be applicable in a wide variety of other domains as well.

Note: It is not specifically built for Deep Learning

# What can TensorFlow do?

---

Linear Models: Linear Regression and Logistics Regression

Neural Network: Shallow NN and Deep Learning

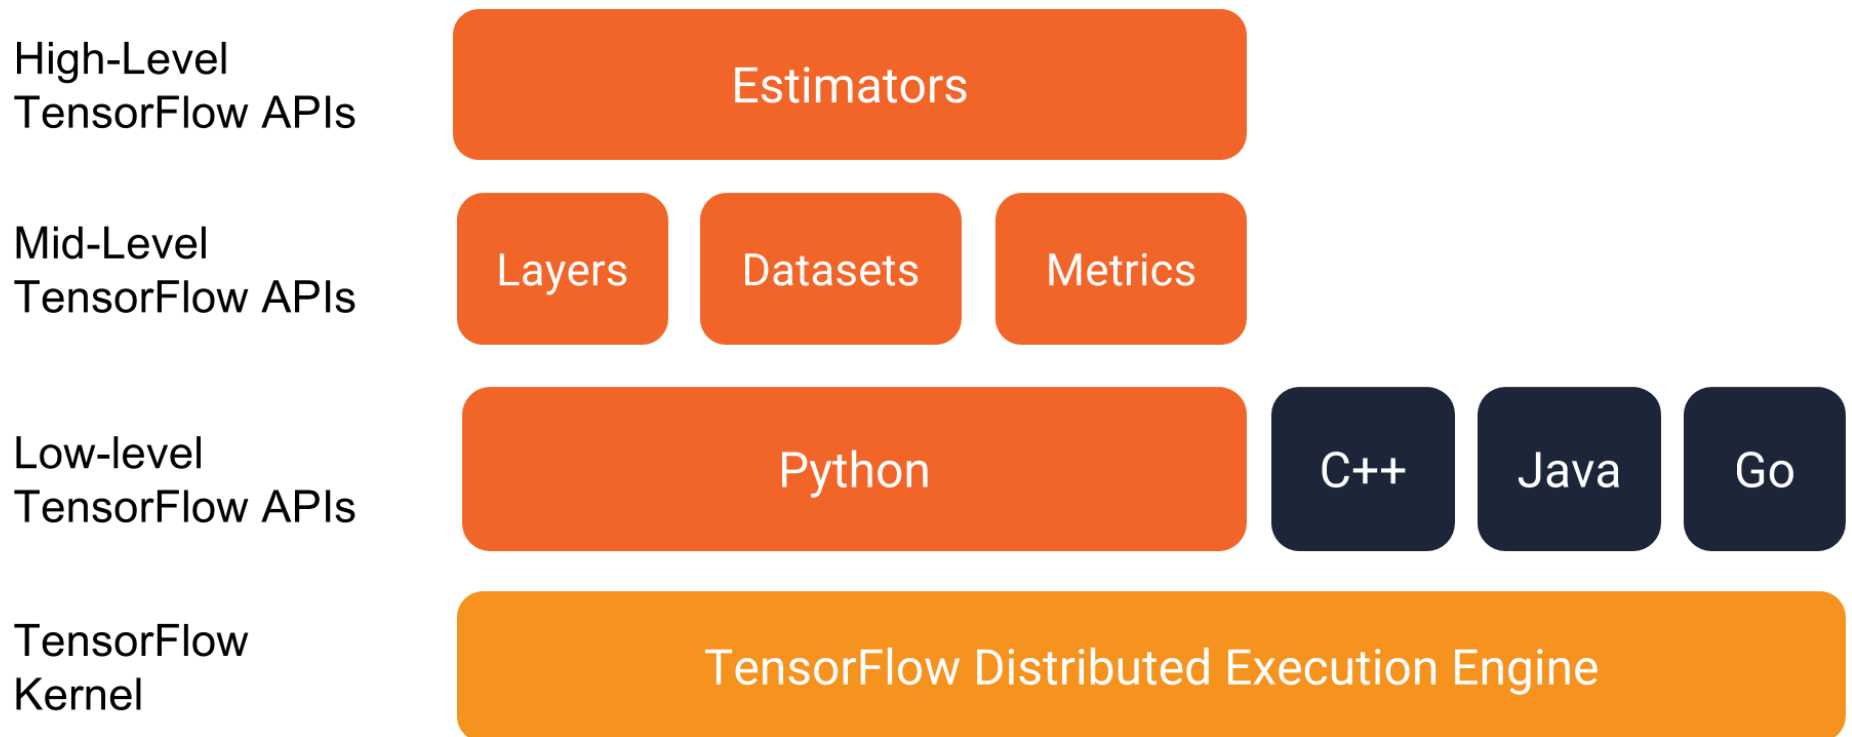
Kernel Methods: Support Vector Machines (SVM)

Other Mathematical Computations: Partial Differential Equations (PDEs)



# TensorFlow Programming

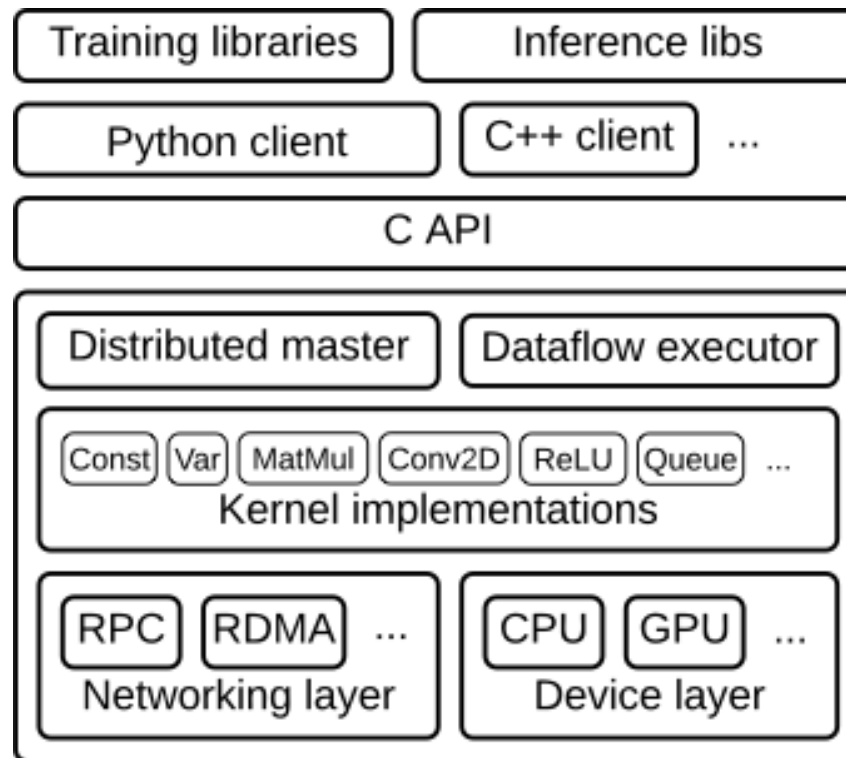
---



source [https://www.tensorflow.org/get\\_started/get\\_started\\_for\\_beginners](https://www.tensorflow.org/get_started/get_started_for_beginners)

# TensorFlow (Runtime stack)

---

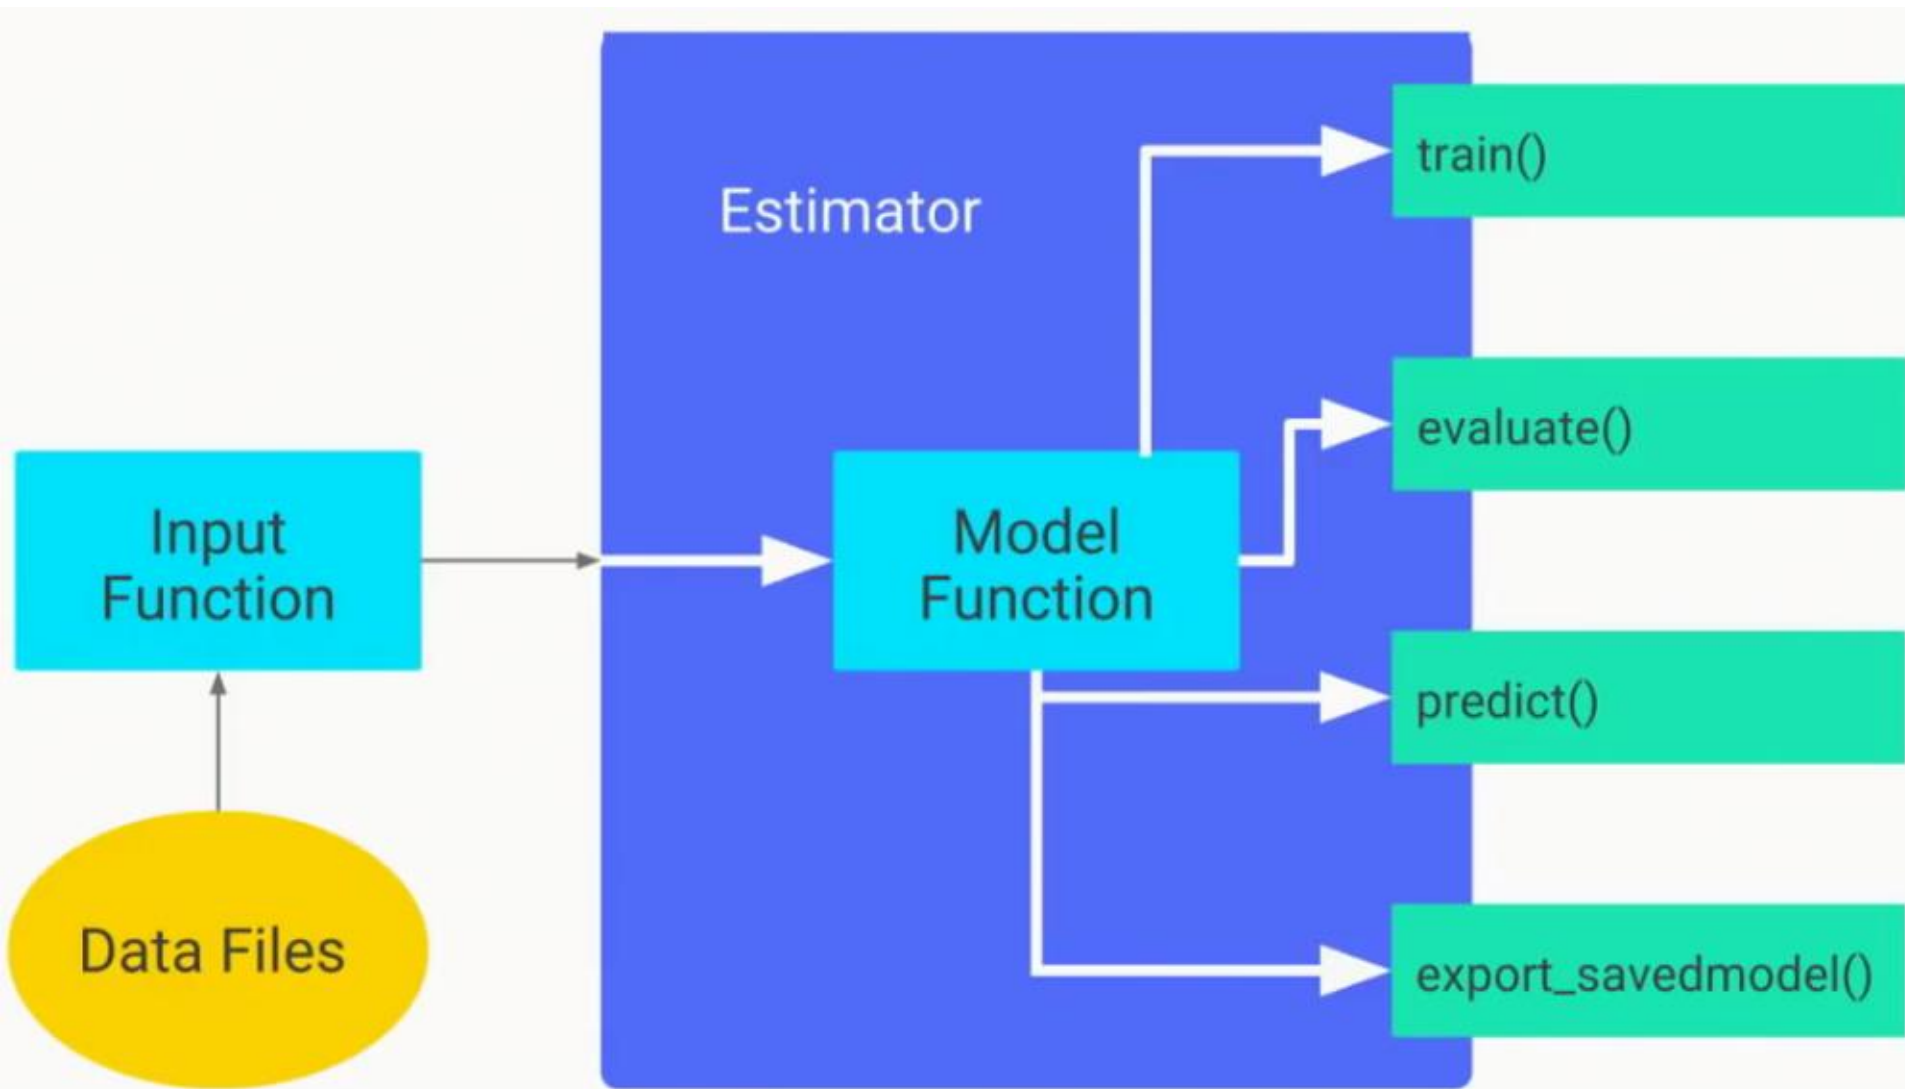


source <https://www.tensorflow.org/extend/architecture>

# TensorFlow API

---

- ❑ TensorFlow supports several programming languages: Python, C++, Java, Go.
- ❑ TensorFlow has APIs available in several languages both for constructing and executing a TensorFlow graph. The Python API is at present the most complete and the easiest to use, but other language APIs may be easier to integrate into projects and may offer some performance advantages in graph execution.
- ❑ Python API: [https://www.tensorflow.org/api\\_guides/python/](https://www.tensorflow.org/api_guides/python/)



# Keras.io

---

- ❑ **Keras** is a high-level neural networks API, written in Python and capable of running on top of **TensorFlow**, **CNTK**, or **Theano**. It was developed with a focus on enabling fast experimentation.
- ❑ Use Keras if you need a deep learning library that:
  - Allows for easy and fast prototyping (through user friendliness, modularity, and extensibility).
  - Supports both convolutional networks and recurrent networks, as well as combinations of the two.
  - Runs seamlessly on CPU and GPU.

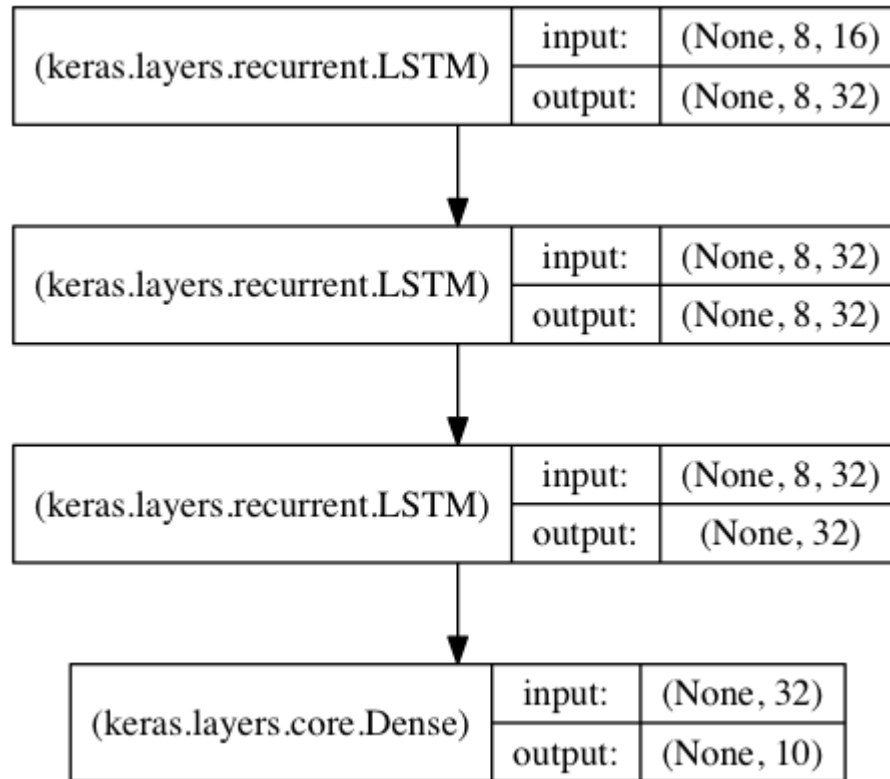
# Keras

---

- ❑ The core data structure of Keras is a **model**, a way to organize layers.
- ❑ The simplest type of model is the **Sequential model**, a linear stack of layers.
- ❑ For more complex architectures, you should use the Keras **functional API**, which allows to build arbitrary graphs of layers.

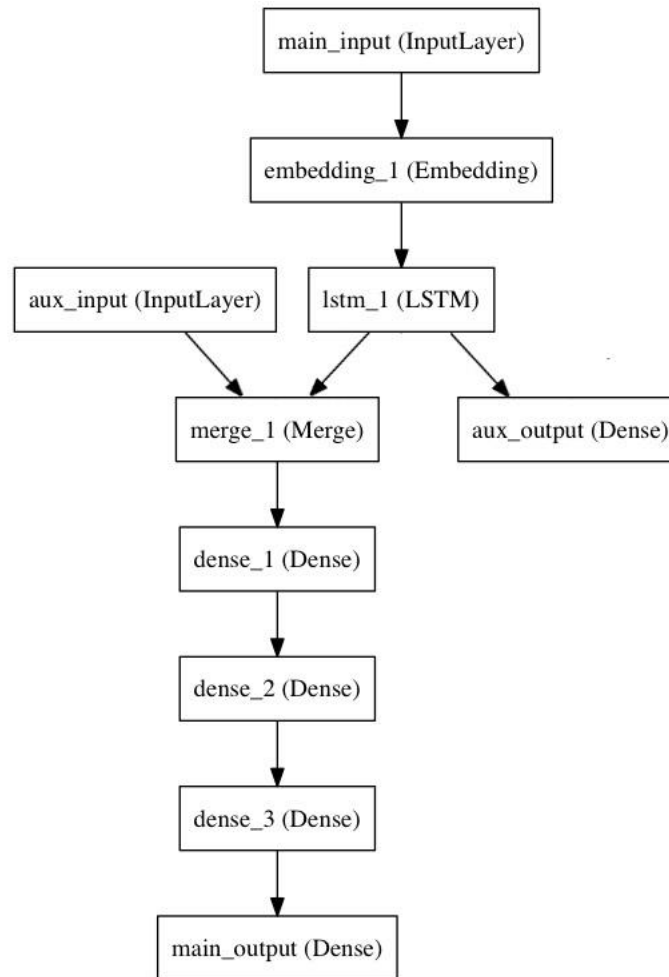
## Stacked LSTM for sequence classification

The Sequential model is a linear stack of layers



source <https://keras.io/getting-started/sequential-model-guide/>

The Keras **functional API** is the way to go for defining complex models, such as multi-output models, directed acyclic graphs, or models with shared layers.







# Distributed Training

# Distributed Training with Tensorflow

---

Source: [https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)

- ❑ `tf.distribute.Strategy` is a TensorFlow API to distribute training across multiple GPUs, multiple machines or TPUs. Using this API, you can distribute your existing models and training code with minimal code changes.
- ❑ `tf.distribute.Strategy` has been designed with these key goals in mind:
  - Easy to use and support multiple user segments, including researchers, ML engineers, etc.
  - Provide good performance out of the box.
  - Easy switching between strategies.
- ❑ `tf.distribute.Strategy` can be used with a high-level API like Keras, and can also be used to distribute custom training loops (and, in general, any computation using TensorFlow).

# Distributed Training with Keras

---

Source:

[https://www.tensorflow.org/guide/distributed\\_training#using\\_tfdistributestrategy\\_with\\_tfkerasmodelfit](https://www.tensorflow.org/guide/distributed_training#using_tfdistributestrategy_with_tfkerasmodelfit)

`tf.distribute.Strategy` is integrated into `tf.keras` which is TensorFlow's implementation of the Keras API specification. `tf.keras` is a high-level API to build and train models. By integrating into `tf.keras` backend, we've made it seamless for you to distribute your training written in the Keras training framework using `model.fit`.

```
mirrored_strategy = tf.distribute.MirroredStrategy()

with mirrored_strategy.scope():
    model = tf.keras.Sequential([tf.keras.layers.Dense(1, input_shape=(1,))])

model.compile(loss='mse', optimizer='sgd')
```

# Distributed Training with Keras

---

Source:

[https://www.tensorflow.org/guide/distributed\\_training#using\\_tfdistributedstrategy\\_with\\_tfkerasmodelfit](https://www.tensorflow.org/guide/distributed_training#using_tfdistributedstrategy_with_tfkerasmodelfit)

This example uses `MirroredStrategy` so you can run this on a machine with multiple GPUs. `strategy.scope()` indicates to Keras which strategy to use to distribute the training. Creating models/optimizers/metrics inside this scope allows us to create distributed variables instead of regular variables. Once this is set up, you can fit your model like you would normally. `MirroredStrategy` takes care of replicating the model's training on the available GPUs, aggregating gradients, and more.

Source: <https://www.tensorflow.org/tutorials/distribute/keras>

Source: [https://www.tensorflow.org/tutorials/distribute/multi\\_worker\\_with\\_keras](https://www.tensorflow.org/tutorials/distribute/multi_worker_with_keras)

# Horovod

---

Source: <https://horovod.ai/>

Horovod is a distributed deep learning training framework for TensorFlow, Keras, PyTorch, and Apache MXNet

Source: <https://github.com/horovod/horovod>

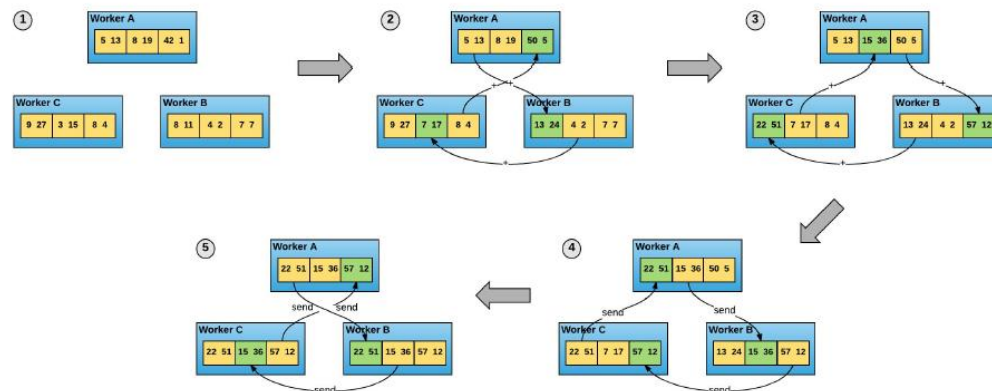
Source: <https://horovod.readthedocs.io/en/stable/>

# Distributed DL with Horovod

Source:

<https://developer.download.nvidia.com/video/gputechconf/gtc/2019/presentation/s9321-distributed-deep-learning-with-horovod.pdf>

Horovod Technique: Ring-Allreduce

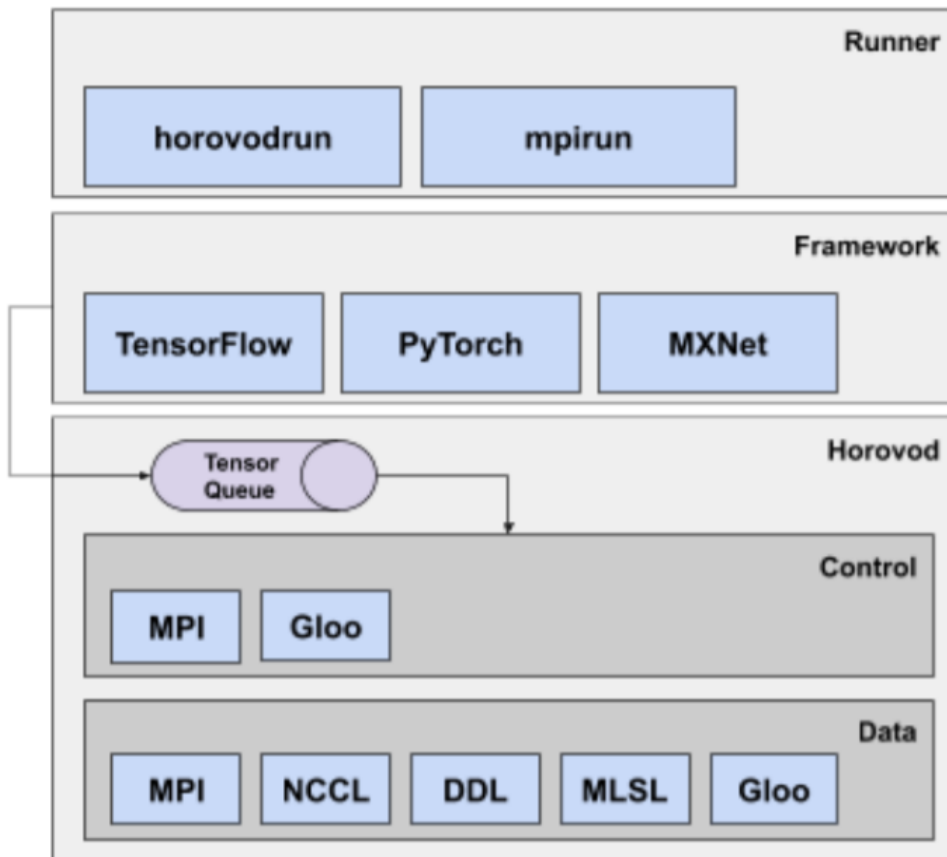


Patarasuk, P., & Yuan, X. (2009). Bandwidth optimal all-reduce algorithms for clusters of workstations. *Journal of Parallel and Distributed Computing*, 69(2), 117-124. doi:10.1016/j.jpdc.2008.09.002

Source: <https://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.141.5642&rep=rep1&type=pdf>

Source: <https://www.adaltas.com/en/2019/11/15/avoid-deep-learning-bottlenecks/>

# Horovod Stack





# Model Zoo

A REPOSITORY OF PRE-TRAINED MODELS FOR A PARTICULAR DEEP LEARNING FRAMEWORK



# Tensorflow Model Zoo

---

□ <https://github.com/tensorflow/models>

- This repository contains a number of different models implemented in TensorFlow:
  - The official models are a collection of example models that use TensorFlow's high-level APIs. They are intended to be well-maintained, tested, and kept up to date with the latest stable TensorFlow API. They should also be reasonably optimized for fast performance while still being easy to read. We especially recommend newer TensorFlow users to start here.
  - The research models are a large collection of models implemented in TensorFlow by researchers. They are not officially supported or available in release branches; it is up to the individual researchers to maintain the models and/or provide support on issues and pull requests.
  - The samples folder contains code snippets and smaller models that demonstrate features of TensorFlow, including code presented in various blog posts.

# Keras Model Zoo

---

❑ <https://keras.io/applications/#available-models>

❑ **Keras Applications** are **deep learning models** that are made available alongside **pre-trained weights**. These models can be used for prediction, feature extraction, and fine-tuning..

❑ Weights are downloaded automatically when instantiating a model. They are stored at `~/.keras/models/`.

❑ Models for image classification with weights trained on ImageNet:

- Xception
- VGG16
- VGG19
- ResNet50
- InceptionV3
- InceptionResNetV2
- MobileNet
- DenseNet
- NASNet


Source <https://github.com/fchollet/deep-learning-models>

# Review of Deep Learning Architectures

---

# Review: LeNet-5

[LeCun et al., 1998]




Conv filters were 5x5, applied at stride 1  
Subsampling (Pooling) layers were 2x2 applied at stride 2  
i.e. architecture is [CONV-POOL-CONV-POOL-FC-FC]

# Case Study: AlexNet

[Krizhevsky et al. 2012]

## Architecture:

CONV1  
MAX POOL1  
NORM1  
CONV2  
MAX POOL2  
NORM2  
CONV3  
CONV4  
CONV5  
Max POOL3  
FC6  
FC7  
FC8



# Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks


8 layers (AlexNet)

-> 16 - 19 layers (VGG16Net)

Only 3x3 CONV stride 1, pad 1  
and 2x2 MAX POOL stride 2

11.7% top 5 error in ILSVRC'13  
(ZFNet)

-> 7.3% top 5 error in ILSVRC'14



AlexNet



VGG16



VGG19

# Case Study: GoogLeNet

[Szegedy et al., 2014]

Deeper networks, with computational efficiency


- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!  
12x less than AlexNet
- ILSVRC’14 classification winner  
(6.7% top 5 error)




# Case Study: GoogLeNet

[Szegedy et al., 2014]

“Inception module”: design a good local network topology (network within a network) and then stack these modules on top of each other



Inception module







# Case Study: ResNet


[He et al., 2015]

Very deep networks using residual connections


- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!




# ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



## Comparing complexity...



ResNet:  
Moderate efficiency depending on  
model, highest accuracy



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

# Summary

---

What we have learnt:

- Tensorflow and Keras can be used to implement Deep Learning
- VGG, GoogLeNet, ResNet all in wide use, available in model zoos
- Trend towards extremely deep networks
- Significant research centers around design of layer / skip connections and improving gradient flow